

ODRF: An R Package for Oblique Decision Tree and Its Random Forest

by Yu Liu and Yingcun Xia

Abstract The classification and regression tree (CART) and Random Forest (RF) are popular machine learning methods that involve selecting one predictor at a time as the splitting variable for each node. The use of linear combinations of predictors as splitting variables is one of the popular extensions of CART and RF, known as Oblique Decision Trees (ODT), and hence ODT-based Random Forests (ODRF). Recent studies have also shown the theoretical advantages of ODT and ODRF over CART and RF. However, there is still no integrated and efficient software package that can demonstrate the numerical advantages of ODT and ODRF. To fill this gap, we make several modifications to the existing algorithms and develop an R package **ODRF** for both `ODT()` and `ODRF()`. In addition, the R package includes a series of new functions, such as ODT-based boosting trees (BODT), the ensemble version of BODT (ODBT), ODT visualization (`plot.ODT`), and an online update function for ODT (`online.ODT`). The main part of ODRF is executed using the **Rcpp** package. This paper presents the basic idea of our modifications and illustrates the use of the package. Through numerical experiments, ODRF was compared with other packages of decision trees and forests, showing a clear overall improvement in classification and regression tasks.

1 Introduction

The Classification and Regression Tree (CART) proposed by Professor Leo Breiman (Breiman et al., 1984) has attracted a great deal of attention from statisticians and data analysts of other disciplines. The method is widely used because it is easy to train and the resulting tree makes the analysis results easy to visualize and interpret (Johnson and Tong, 2014). On the other hand, much attention has been paid to the algorithm and many improvements have been proposed. Classification and regression trees (Quinlan, 1987, CART) and C4.5 (Quinlan, 1993) are perhaps the most commonly used decision trees. There is a long list of other decision trees, including the Evolutionary Learning of Globally Optimal Classification and Regression Trees (Grubinger et al., 2014, EVT), Conditional Inference Trees (Hothorn et al., 2006, CT), Extremely randomized trees (Geurts et al., 2006, ERT), Model-Based Recursive Partitioning (Zeileis and Hothorn, 2015, MOB), Bayesian additive regression trees (Maia et al., 2022, (BART)) and generalized linear mixed-model trees (Fokkema et al., 2020, glmertree). The Random Forests (Breiman, 2001, RF), which is an ensemble of trees by either feature bagging or boosting, is arguably one of the most efficient machine learning methods, especially for the tabular data (Grinsztajn et al., 2022). Again, there are many ensemble methods that are based on different decision trees, for example, Conditional Random Forests (Hothorn et al., 2006, cforest), Learning Nonlinear Functions Using Regularized Greedy Forest (Johnson and Tong, 2014, RGF) and Generalized Random Forest (Athey et al., 2019, GRF).

One of the most appealing extensions to CART is the use of linear combinations of the predictors as splitting variables that is known as the Oblique Decision Tree (Heath et al., 1993, ODT). Recently, Zhan et al. (2025) proved the consistency of ODT and its random forest (ODRF) for very general regression functions as long as they are continuous, while CART or RF are consistent mainly for regressions with special structures such as additive models. Again, ensemble can be made based on ODT, resulting in the Oblique-type Random Forests, including Random Rotation Random Forest (RotRF) of Blaser and Fryzlewicz (2016), Random Projection Forests (RPFs) of Lee et al. (2015) and Sparse Projection Oblique Randomer Forests (SPORF) of Tomita et al. (2020). Another type of random forests is the model-based oblique decision forest, including mainly Canonical Correlation Forests (CCF) with classic correlation analysis (Rainforth and Wood, 2015), projection pursuit forest (PPF) with linear discriminant analysis (da Silva et al., 2021), oblique random forests (ORF) with ridge regression (Menze et al., 2011), oblique random survival forests (Jaeger et al., 2022, ORSF) and Heterogeneous oblique random forest (Katuwal et al., 2020, HORF). However, a primary limitation of these oblique methods lies in their computational intensity. To address this issue, Wickramarachchi et al. (2016) introduced a novel algorithm named “HHCART”, which utilizes Householder matrices to reflect training data at non-terminal nodes. Subsequently, Wickramarachchi et al. (2019) developed an accompanying R package **hhcart** for both classification and regression tasks.

In parallel with these developments, Gradient Boosting Trees (GBT) have emerged as another powerful machine learning paradigm. The GBT algorithm was first introduced by Friedman (2001), establishing a powerful framework for supervised learning by sequentially fitting weak learners (typically decision trees) to the residuals of previous models. To address overfitting, Friedman (2002)

later proposed stochastic gradient boosting, which incorporates regularization by subsampling both observations and features. Further enhancements tailored to tree-based models were introduced through modern implementations such as XGBoost (Chen and Guestrin, 2016), LightGBM (Ke et al., 2017), and CatBoost (Prokhorenkova et al., 2018). These frameworks optimize tree growth strategies, efficiently handle categorical variables, and significantly improve computational scalability. Traditionally, boosting algorithms employed axis-aligned splits (as in CART-style trees). However, recent research has explored Oblique Decision Trees (ODT) within boosting frameworks, resulting in Boosting Oblique Decision Trees (BODT). Integrating BODT with randomization techniques, such as random projections (Lee et al., 2015), further enhances robustness and generalizability. Notably, our recent work demonstrated that ensembles based on BODT achieve consistency for general continuous regression functions, in contrast to traditional CART-based ensembles that require additive or structured assumptions.

Although the theoretical advantages of oblique decision trees and their random forests have been well understood, the existing packages implementing those extensions only show their better numerical performance in some special cases, and thus have not commonly received and have not got much popularity as the methods deserve. As a consequence, the conventional packages **rpart** (Therneau and Atkinson, 2022) of Therneau and Atkinson (2000) for CART and **randomForest** (Liaw and Wiener, 2002a) of Liaw and Wiener (2002b) for RF are still the most commonly used packages. The main difficulty in implementing ODT or ODRF is the estimation of the coefficient, θ , for the linear combinations, which is also one of the main differences amongst all the existing packages. The estimation methods of θ include random projection, logistic regression, dimension reduction and many others. For example, package **rerf** (Browne and Tomita, 2019) of Tomita et al. (2020) uses random projections; package **PPforest** (da Silva et al., 2025) of da Silva et al. (2021) uses linear discriminant and penalized discriminant analysis; package **obliqueRF** (Menze and Splitthoff, 2012) of Menze et al. (2011) uses the partial least squares regression, logistic regression, and random projection to find the projections. On the other hand, most of the existing R packages for such types of random forests can only be used for classification. Some of the R packages have been removed from the Comprehensive R Archive Network (CRAN) at <https://CRAN.R-project.org/>. For comprehensive details, see Table 1, which presents a comparative analysis of nine R packages implementing oblique tree methods. The evaluation assesses their support for six specific functionalities and determines whether they remain available on CRAN.

Our package, called **ODRF**, is written based on the above-mentioned packages and the recent work of Zhan et al. (2025). In **ODRF**, the projection pursuit regression (Friedman and Stuetzle, 1981) is used to find θ ; the details will be stated in the following section. Other options for the estimations of projections are also provided in the package. Compared with the existing random forests, the advantages of **ODRF** are as follows.

- **ODRF** can be used for both classification and regression, while most existing packages of oblique-type trees or forests can only be used for classification.
- Both the tree (ODT) and random forests (ODRF) of **ODRF** have better overall accuracy than existing trees and forests, including traditional CART and RF and other oblique trees or forests.
- **ODRF** allows users to define their own functions to find the projections at each node, which is essential to the performance of the forests.
- **ODRF** provides an online class of functions that facilitate the application to streaming data, enabling continuous updates and improvements to existing trees or random forests.
- **ODRF** develops a new ensemble method (ODBT), which builds random forests of ODT-based boosting trees and outperforms traditional boosting trees and random forests in both classification and regression tasks.
- **ODRF** proposes an innovative approach for constructing linear model trees by incorporating LASSO regularization, while distinctly separating the sets of variables used for node splitting and model fitting.

The remainder of this paper is organized as follows: Section 2 describes the model or methods of the main functions in **ODRF**. Section 3 provides the usage of the main functions in **ODRF** package. Section 4 showcases the specific application of **ODRF** to analyze two data sets with continuous and categorical responses respectively, and compares the predictive accuracy in classification and regression with other R packages using 43 datasets. A conclusion is made in Section 5.

2 Statistical methods

Suppose $Y = (y_1, \dots, y_K)$ is the response vector of interest and $X = (x_1, \dots, x_p)^\top : p \times 1$ is the vector of predictors. We allow Y to be multiple to accommodate the categorical response. That is, if Y has K

	classification	regression	forests	boosting	custom splits	online updating	CRAN
oblique.tree	✓	✗	✗	✗	✗	✗	✗
pptreeviz	✓	✗	✗	✗	✗	✗	✓
PPtreereg	✗	✓	✗	✗	✗	✗	✗
rotationForest	✓	✗	✓	✗	✗	✗	✓
rerf	✓	✗	✓	✗	✗	✗	✗
PPforest	✓	✗	✓	✗	✗	✗	✓
obliqueRF	✓	✗	✓	✗	✗	✗	✗
aorsf	✓	✓	✓	✗	✗	✗	✓
ODRF	✓	✓	✓	✓	✓	✓	✓

Table 1: A comprehensive overview of the R package for oblique tree methods

classes, then it is represented by K dummy variables, with each taking values 0 and 1. Generally, we need to estimate the regression functions:

$$m(x) = (m_1(x), \dots, m_K(x)) = (E(y_1|X = x), \dots, E(y_K|X = x)).$$

2.1 Create an ODT

With observations $\mathbb{A}_0^0 = \{(X_i, Y_i), i = 1, \dots, n\}$, where $Y_i = (y_{i1}, \dots, y_{iK})$ and $X_i = (x_{i1}, \dots, x_{ip})$, an ODT is illustrated by the following diagram in Figure 1. For ease of exposition, even if a node is not split further, we still rewrite it in the next layer (by a dashed line in the diagram). For any node $\mathbb{A}_\ell^\tau$, where the subscript ℓ represents the layer and superscript the number of nodes in the layer, the splitting is as follows. Given any p -dimensional vector θ and a splitting value c , define the daughter nodes by

$$\mathbb{A}_{\ell+1}^{\tau'} = \{X_i : X_i \in \mathbb{A}_\ell^\tau, \theta^\top X_i \leq c\}, \quad \mathbb{A}_{\ell+1}^{\tau''} = \{X_i : X_i \in \mathbb{A}_\ell^\tau, \theta^\top X_i > c\},$$

and loss function

$$\Delta(c|\mathbb{A}_\ell^\tau, \theta) = \sum_{k=1}^K \sum_{X_i \in \mathbb{A}_{\ell+1}^{\tau'}} (y_{ik} - \bar{y}_k(\mathbb{A}_{\ell+1}^{\tau'}))^2 + \sum_{k=1}^K \sum_{X_i \in \mathbb{A}_{\ell+1}^{\tau''}} (y_{ik} - \bar{y}_k(\mathbb{A}_{\ell+1}^{\tau''}))^2,$$

where $\bar{y}_k(\mathbb{A}) = \sum_{X_i \in \mathbb{A}} y_{ik} / \#\mathbb{A}$ and $\#\mathbb{A}$ denotes the cardinality of set $\mathbb{A}$. When θ is given, the splitting value c should minimize $\Delta(c|\theta, \mathbb{A}_\ell^\tau)$. If Y is a univariate quantitative response, then $\Delta(c|\theta, \mathbb{A}_\ell^\tau)$ is the residual sum of squares for regression that is used as the criterion of CART. If Y is categorical, then $\bar{y}_k(\mathbb{A}) = \hat{p}_k(\mathbb{A})$, the ratio of category k in $\mathbb{A}_{\ell+1}^{\tau'}$, and

$$\sum_{X_i \in \mathbb{A}} (y_{ik} - \bar{y}_k(\mathbb{A}))^2 = \#\mathbb{A} \times (1 - \hat{p}_k(\mathbb{A}))\hat{p}_k(\mathbb{A}).$$

Thus, $\Delta(c|\theta, \mathbb{A}_\ell^\tau)$ is also the Gini impurity but multiplied by the number of observations in the node, which thus is also the criterion used by CART for classification.

To determine whether a split is necessary, the cross-validation (CV) method is used as follows. For any set $\mathbb{A}$, the CV value is defined as a scaled residual sum of squares:

$$CV(\mathbb{A}) = \left(\frac{\#\mathbb{A}}{\#\mathbb{A} - \lambda} \right)^2 \sum_{k=1}^K \sum_{X_i \in \mathbb{A}} (y_{ik} - \bar{y}_k(\mathbb{A}))^2,$$

Note that if $\lambda = 1$, then $CV(\mathbb{A})$ is the leave-one-out CV. We can show that $\lambda = \log(\#\mathbb{A})$ also gives a consistent stopping time. In our package, if

$$CV(\mathbb{A}_\ell^\tau) \leq CV(\mathbb{A}_{\ell+1}^{\tau'}) + CV(\mathbb{A}_{\ell+1}^{\tau''}),$$

then $\mathbb{A}_\ell^\tau$ is a leaf and no longer needs to be split; otherwise, it will be split to daughter nodes $\mathbb{A}_{\ell+1}^{\tau'}$ and $\mathbb{A}_{\ell+1}^{\tau''}$; see the first equation of this section.

Let $\{\mathbb{A}_L^j\}_{j=1}^{t_n}$ be all the leaves (i.e. the nodes in the last layer) of the generated tree. Then, $m(x)$ is




Figure 1: A diagram of the oblique decision tree

estimated by

$$m_n(x) = \sum_{j=1}^{t_n} \mathbb{I}(x \in \mathbb{A}_L^j) \times (\bar{y}_1(\mathbb{A}_L^j), \dots, \bar{y}_K(\mathbb{A}_L^j)).$$

Note that if Y is categorical, $m_n(x)$ is the vector of probabilities of each class.

It can be seen that the main step in ODT is the estimation of the projection θ , i.e. the coefficients of the linear combinations. Although many methods have been proposed as mentioned in the Introduction section, we find the projection pursuit regression (Friedman and Stuetzle, 1981) is still the most efficient and is used in ODRF. The estimation is as follows. In any node $\mathbb{A}_\ell^\tau$, define a loss function

$$\Delta(\theta) = \sum_{k=1}^K \sum_{X_i \in \mathbb{A}_\ell^\tau} \left\{ y_{ik} - m_k(\theta^\top X_i) \right\}^2,$$

where m_k is a nonparametric smoothing of the regression function that minimizes $\sum_{X_i \in \mathbb{A}_\ell^\tau} \left\{ y_{ik} - m_k(\theta^\top X_i) \right\}^2$ with θ given. The nonparametric smoothing can be either a spline representation, kernel smoothing, or a super smoother of Friedman (1984). Projection θ is estimated as

$$\theta = \arg \min_{\theta} \Delta(\theta).$$

Our package also allows users to define their own methods of estimating θ ; see Section 3.5.

2.2 Build an ODRF

ODRF builds the random forest slightly differently from the existing forests, but its basic idea is still the feature bagging (Ho, 1998). The detail is as follows. Denote by $X_{[q]} = (x'_1, \dots, x'_q)$ a random subset of predictors $X = (x_1, \dots, x_p)$, where $q < p$ and $\{x'_1, \dots, x'_q\} \subset \{x_1, \dots, x_p\}$, and thus by $X_{[q],r}$, $r = 1, 2, \dots$, we mean a sequence of such subsets that may differ from one another. In other words, with the same q and r , set $X_{[q],r}$ changes from place to place but has the same cardinality. The ODRF is implemented as follows.

- Create a random ODT tree, labeled as b , using the idea of feature bagging as follows.
 - For each node $\mathbb{A}_k^\tau$, randomly select q , for example $q = \lfloor p/3 \rfloor$ or $q = \lfloor \sqrt{p} \rfloor$, variables $X_{[q],r} = (x'_{11}, \dots, x'_{iq}) \subset X$, where $r = 1, \dots, R$ denotes R random sets of variables. In ODRF, $R \geq \lfloor p/q \rfloor$ is used as default.
 - For each $X_{[q],r}$, find

$$\tilde{\theta}_r = \arg \min_{\theta} \sum_{k=1}^K \sum_{X_i \in \mathbb{A}_k^\tau} \left\{ y_{ik} - m_k(\theta^\top X_{[q],r,i}) \right\}^2,$$

$$\text{where } X_{[q],r,i} = (x'_{i1}, \dots, x'_{iq})^\top.$$

- Define

$$\mathbb{A}_{k,r}^{\tau'} = \{X_i : X_i \in \mathbb{A}_k^{\tau}, \theta_{(q)}^{\top} X_{[q],r,i} \leq c_{(q)}\}, \quad \mathbb{A}_{k,r}^{\tau''} = \{X_i : X_i \in \mathbb{A}_k^{\tau}, \theta_{(q)}^{\top} X_{[q],r,i} > c_{(q)}\}.$$

- Calculate

$$(\tilde{c}_{(q)}, \tilde{r}) = \arg \min_{c,r=1,\dots,R} \left\{ \sum_{k=1}^K \sum_{X_i \in \mathbb{A}_{\ell,r}^{\tau'}} (y_{ik} - \bar{y}_{.k})^2 + \sum_{k=1}^K \sum_{X_i \in \mathbb{A}_{\ell,r}^{\tau''}} (y_{ik} - \bar{y}_{.k})^2 \right\}.$$

- Split the node with $(\tilde{\theta}_{\tilde{r}}, \tilde{c}_{\tilde{r}}, \tilde{r})$ to produce daughter nodes

$$\mathbb{A}_{\ell+1}^{\tau'} = \mathbb{A}_{\ell+1,\tilde{r}}^{\tau'} \text{ and } \mathbb{A}_{\ell+1}^{\tau''} = \mathbb{A}_{\ell+1,\tilde{r}}^{\tau''}.$$

If $CV(\mathbb{A}_k^{\tau}) < CV(\mathbb{A}_{\ell+1}^{\tau'}) + CV(\mathbb{A}_{\ell+1}^{\tau''})$, we discard the two daughter nodes, and make $\mathbb{A}_k^{\tau}$ as a leaf; otherwise, we continue to split the nodes.

- Each tree produces one estimator $\hat{m}^b(x) = (\hat{m}_1^b(x), \dots, \hat{m}_K^b(x))$.

- An Oblique Decision Random forest (ODRF) estimator is then

$$\hat{m}_{ODRF}(x) = B^{-1} \sum_{b=1}^B \hat{m}_n^b(x).$$

2.3 Build an ODBT

Building upon the foundational gradient boosting framework established by [Friedman \(2001\)](#), the boosting process constructs trees sequentially. Specifically, the k -th tree is trained using the predictors and residuals derived from the previous $k-1$ trees. Consequently, the estimator $m^{(k)}(x)$ at step k is expressed as a linear combination of the first k trees. In this work, we propose using Oblique Decision Trees (ODT) as base learners within the boosting framework. The resulting ODT-based boosting algorithm (BODT) proceeds according to the following key steps:

- Step 1. Initialize $k = 1$, training set $\mathcal{D}_n^1 = \mathbb{A}_0^0$, residuals $r_{0,i} = Y_i$ for $i = 1, \dots, n$, and estimator $m_{t,0}^r = 0$.
- Step 2. Train the k -th ODT using data $\mathcal{D}_n^k = \{(X_i, r_{k-1,i})\}_{i=1}^n$. Denote the resulting tree as $m_{t,k}^r(x)$. (Note that for $k = 1$, $m_{t,1}^r$ is simply the ODT trained on $\mathcal{D}_n$.)
- Step 3. Update the estimator of $m(x)$ at step k by:

$$m_{k,\text{boost}}(x) := \sum_{j=1}^k a_{k,j}^* m_{t,j}^r(x), \quad (2.1)$$

where the coefficients $(a_{k,1}^*, a_{k,2}^*, \dots, a_{k,k}^*)$ are obtained by solving:

$$(a_{k,1}^*, \dots, a_{k,k}^*) := \arg \min_{(a_{k,1}, \dots, a_{k,k}) \in \mathbb{R}^k} \sum_{i=1}^n \left(Y_i - \sum_{j=1}^k a_{k,j} m_{t,j}^r(X_i) \right)^2. \quad (2.2)$$

- Step 4. Update the residuals as $r_{k,i} = Y_i - m_{k,\text{boost}}(X_i)$ and the training data as $\mathcal{D}_n^{k+1} = \{(X_i, r_{k,i})\}_{i=1}^n$.
- Step 5. Set $k \leftarrow k + 1$ and return to Step 2.

To further improve the stability and predictive accuracy of the boosting trees, we incorporate a bagging step by replacing each deterministic ODT with a randomly constructed ODT. For simplicity, we set $a_n = n$, meaning the full dataset $\mathcal{D}_n$ is used at every boosting iteration. During the construction of each random ODT, independent random seeds (as defined at the beginning of Section 2.2) are employed. By a slight abuse of notation, we continue to denote the resulting estimator as $m_{k,\text{boost}}(x)$. Running the BODT procedure independently B times, each with its own set of random ODTs, produces an ensemble estimator, referred to as ODBT, defined by

$$m_{k,\text{boost}}^{\text{ens}}(x) := \frac{1}{B} \sum_{b=1}^B m_{k,\text{boost}}^b(x),$$

where $m_{k,\text{boost}}^b(x)$ corresponds to the boosting estimator from the b -th independent run.

For a comprehensive discussion of the technical details and theoretical properties of ODBT, we refer the reader to our parallel work.

2.4 Online training with batches of data

The package allows for easy model updating (for ODT or ODRF) when new data is available. Essentially, we achieve this by splitting a leaf node upon the arrival of new data. The specifics of this process vary based on the amount of available data. To keep things simple, we focus on classification trees that use Gini impurity as the splitting criterion.

Suppose the trained ODT has leaves $A_L^j, j = 1, \dots, J$. When the new data comes, we can simply fit the new data into the leaves of the trained tree. Suppose $D_L^j = \{(X'_i, y'_i), i = 1, \dots, n'_j\}$ fall into A_L^j , where $y'_i = (y'_{i,1}, \dots, y'_{i,K})$, thus the data in the leaf becomes $A_L^j \cup D_L^j$. Next, we discuss whether we need to split $A_L^j \cup D_L^j$ in different scenarios of data availability.

- If the full data for the original trained model is known, that is, the data in $A_L^j \cup D_L^j$ is fully known, we can use the method of ODT to split the data.
- If only the impurity is known, i.e. $n_j = \#A_L^j$ and $(\hat{p}_{L,1}^j, \dots, \hat{p}_{L,K}^j)$ are kept for the leaf, which is indeed the case for many existing packages of trees. Here, $\hat{p}_{L,k}^j$ is the probability of class k in the leaf A_L^j . The impurity of A_L^j and $A_L^j \cup D_L^j$ can be calculated respectively as

$$G_L^j = \sum_{k=1}^K \hat{p}_k(1 - \hat{p}_k)$$

and

$$\tilde{G}_L^j = \sum_{k=1}^K \tilde{p}_k(1 - \tilde{p}_k),$$

where

$$\tilde{p}_k = (n_j \hat{p}_{L,k}^j + n'_j \check{p}_{L,k}^j) / (n_j + n'_j)$$

and $n'_j = \#D_L^j$ and $\check{p}_{L,k}^j = \sum_{X'_i \in D_L^j} y'_{i,k} / n'_j$. If $\hat{G}_L^j \leq \tilde{G}_L^j$, we don't need to split. Otherwise, we first split D_L^j in the same way as ODT does, and denote the daughter nodes by $A_L^{j,l}$ and $A_L^{j,r}$. We assume that the original data in the leaf are also proportionally split into the two daughter nodes. Thus, their impurities are

$$G(A_L^{j,l}) = \sum_{k=1}^K \hat{p}_k^l(1 - \hat{p}_k^l), \quad G(A_L^{j,r}) = \sum_{k=1}^K \hat{p}_k^r(1 - \hat{p}_k^r),$$

where $n_j^l = \tau n_j + \#A_L^{j,l}$ and $n_j^r = (1 - \tau)n_j + \#A_L^{j,r}$, $\hat{p}_k^l = (\tau n_j \hat{p}_{L,k}^j + \sum_{X'_i \in A_L^{j,l}} y'_{i,k}) / n_j^l$ and $\hat{p}_k^r = ((1 - \tau)n_j \hat{p}_{L,k}^j + \sum_{X'_i \in A_L^{j,r}} y'_{i,k}) / n_j^r$, and $\tau = \#A_L^{j,l} / n'_j$. Following the same rule as ODT in the splitting, if

$$\left(\frac{n_j + n'_j}{n_j + n'_j - \lambda} \right)^2 (n_j + n'_j) \tilde{G}_j \leq \left(\frac{n_j^l}{n_j^l - \lambda} \right)^2 n_j^l G(A_L^{j,l}) + \left(\frac{n_j^r}{n_j^r - \lambda} \right)^2 n_j^r G(A_L^{j,r}),$$

then we split the leaf, and continue to split its daughter nodes if necessary; otherwise, A_L^j is not split and remains as a leaf.

The above procedure is a combination of the batches of data in the leaves. In **ODRF**, the second scenario is used as we believe most trees have the data required.

3 Overview of ODRF functions

ODRF is written in **Rcpp** (Eddelbuettel and François, 2011) package and functions of R's S3, including the base R functions `print()`, `predict()`, and `plot()` in the base package, the conversion function `as.party()` in package **partykit** (Hothorn and Zeileis, 2015). Next, we introduce the main functions of **ODRF** package, including:

- `ODT()` constructs a single oblique decision tree for classification and regression tasks, where each node split is determined by a linear combination of predictors. Returns an ODT S3 object compatible with `plot()`, `predict()`, and `print()` methods.
- `ODRF()` implements oblique random forest ensembles for classification and regression, extending the random forest framework through ODT-based construction while encompassing standard random forest as a special case. Returns an ODRF S3 object and supports the
- `ODBT()` implements boosted oblique trees by applying feature bagging during ODT-based boosting training to ensemble multiple models. It accepts training and test data sets as input and directly outputs fitted values and predictions.
- `online()` supports incremental learning scenarios and is applicable to ODT and ODRF S3 objects.
- `VarImp()` computes variable importance measures and is compatible with ODT and ODRF S3 objects.

Finally, we demonstrate the implementation of user-defined projection estimation functions through customizable template components, enabling researchers to tailor hyperplane optimization strategies to specific problem domains.

3.1 Print the tree structure of ODT and ODRF

`ODT()`, `ODRF()`, and `ODBT()` are the three main functions of **ODRF** package. They can be used for classification and regression and are similar in usage to packages **rpart** and **randomForest** respectively.

We provide two ways of data input following the S3 methods as follows.

```
## S3 method for class 'formula'
ODT(formula, data = NULL, split = "auto", NodeRotateFun = "RotMatPPO", ...)
ODRF(formula, data = NULL, split = "auto", NodeRotateFun = "RotMatPPO", ...)
ODBT(formula, data = NULL, Xnew = NULL, split = "auto", model = "ODT", ...)

## Default S3 method:
ODT(X, y, split = "auto", NodeRotateFun = "RotMatPPO", ...)
ODRF(X, y, split = "auto", NodeRotateFun = "RotMatPPO", ...)
ODBT(X, y, Xnew, split = "auto", model = "ODT", ...)
```

The formula and data are standard formats in R, so do the remaining arguments such as subset, na.action, and weights. For additional information on setting the values of arguments that are not introduced here, please refer to the documentation for the functions `ODT` and `ODRF` using commands `?ODT`, `?ODRF`, and `?ODBT`. However, in most cases, the default values work well.

`print()` can be used to display the trained tree and structure of each node of ODT in detail. Here is one example.

```
set.seed(38)
data(iris, package = "datasets")
tree <- ODT(Species ~ ., data = iris)
print(tree)

=====
Oblique Classification Tree structure
=====

1) root
node2)# proj1*X < 0.29 -(leaf1 = setosa)
node3) proj1*X >= 0.29
node4)# proj2*X < 0.88 -(leaf2 = versicolor)
node5)# proj2*X >= 0.88 -(leaf3 = virginica)

party.tree <- as.party(tree, data = iris)
print(party.tree)

Model formula:
Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width

Fitted party:
[1] root
```

```
| [2] proj1*X >= 0.29167
| | [3] proj2*X >= 0.88395: virginica (n = 53, err = 5.7%)
| | [4] proj2*X < 0.88395: versicolor (n = 47, err = 0.0%)
| [5] proj1*X < 0.29167: setosa (n = 50, err = 0.0%)
```

```
Number of inner nodes: 2
Number of terminal nodes: 3
```

In addition, we can also use function `print()` to print the model fitted error for ODRF.

```
set.seed(38)
forest <- ODRF(Species ~ ., data = iris, parallel = FALSE)
print(forest)
```

```
Call:
ODRF.formula(formula = Species ~ ., data = data, parallel = FALSE)
Type of oblique decision random forest: classification
Number of trees: 100
OOB estimate of error rate: 5.33\%
Confusion matrix:
```

	setosa	versicolor	virginica	class_error
setosa	50	0	0	0.0000000
versicolor	0	47	5	0.0961537
virginica	0	3	45	0.0625000

Beyond its core functionality, the ODT implementation can also perform linear model trees (LMT) with LASSO regularization as proposed by [Craig et al. \(2024\)](#). This approach differs fundamentally from standard ODT in its data partitioning: the splitting variables and linear modeling variables are trained on distinct datasets. We demonstrate this functionality through the following illustrative case, where we use “Z” as the splitting variable to build a linear model tree for “X” and “y”.

```
set.seed(10); cutpoint <- 50; mu <- rep(0, 100)
X <- matrix(rnorm(100 * 10), 100, 10)
age <- sample(seq(20, 80), 100, replace = TRUE)
height <- sample(seq(50, 200), 100, replace = TRUE)
weight <- sample(seq(5, 150), 100, replace = TRUE)
Z <- cbind(age = age, height = height, weight = weight)
mu[age <= cutpoint] <- X[age <= cutpoint, 1] + X[age <= cutpoint, 2]
mu[age > cutpoint] <- X[age > cutpoint, 1] + X[age > cutpoint, 3]
y <- mu + rnorm(100)
my.tree <- ODT(X = X, y = y, Xsplit = Z, split = "linear", lambda = 0, NodeRotateFun =
"RotMatRF", glmnetParList = list(lambda = 0, family = "gaussian"))
pred <- predict(my.tree, X, Xsplit = Z)
mean((pred - y)^2)
[1] 0.9035932
```

3.2 Classification and regression using functions `ODT()`, `ODRF()`, and `ODBT()`

`predict()` is the standard S3 method to predict new data for various objects of classes. We defined the functions `predict.ODT()` and `predict.ODRF()` to predict `Xnew` for classes `ODT()` and `ODRF()` respectively. The default output of `predict()` is response which is the predicted values of the new data. Use `?predict.ODT` and `?predict.ODRF` to see the detail of the prediction.

The standard usage is as follows.

```
## S3 method for class 'ODT'
predict(object, Xnew, ...)
## S3 method for class 'ODRF'
predict(object, Xnew, type = "response", ...)
```

Examples of classification and regression using `ODT()`, `ODRF()`, and `ODBT()` are as follows.

```
data(body_fat, package = "ODRF")
train <- sample(1:252, 200)
bodyfat_train <- data.frame(body_fat[train, ])
bodyfat_test <- data.frame(body_fat[-train, ])
tree <- ODT(Density ~ ., bodyfat_train, split = "mse")
```

```

pred <- predict(tree, bodyfat_test[, -1])
(e.tree <- mean((pred - bodyfat_test[, 1])^2))
[1] 4.775053e-05

set.seed(12)
data(seeds, package = "ODRF")
train <- sample(1:209, 150)
seeds_train <- data.frame(seeds[train, ])
seeds_test <- data.frame(seeds[-train, ])
forest <- ODRF(varieties_of_wheat ~ ., seeds_train, split = "gini", parallel = FALSE)
pred <- predict(forest, seeds_test[, -8])
(e.forest <- mean(pred != seeds_test[, 8]))
[1] 0.01694915

forest <- ODBT(varieties_of_wheat ~ ., seeds_train, seeds_test[, -8], model = "rpart",
  type = "class", max.terms = 10, parallel = FALSE, NodeRotateFun = "RotMatRF")
pred <- forest$results$prediction
(mean(pred != seeds_test[, 8]))
[1] 0.01694915

```

3.3 Online updating

ODRF provides online training for sequential data with function `online()` and update existing ODT and ODRF using batches of data.

The usage is as follows.

```

## S3 method for class 'ODT' and 'ODRF'
online(obj, X = NULL, y = NULL, ...)

```

In the following example, the training data are available in two batches. The first batch is used to train ODT and ODRF, and the second batch is used to update the trained model.

```

set.seed(17)
index <- sample(nrow(seeds_train), floor(nrow(seeds_train) / 2))
forest1 <- ODRF(varieties_of_wheat ~ ., seeds_train[index, ], split = "gini",
  parallel = FALSE)
pred <- predict(forest1, seeds_test[, -8])
(e.forest.1 <- mean(pred != seeds_test[, 8]))

[1] 0.03389831

forest2 <- online(forest1, seeds_train[-index, -8], seeds_train[-index, 8])
pred <- predict(forest2, seeds_test[, -8])
(e.forest.online <- mean(pred != seeds_test[, 8]))

[1] 0.01694915

index <- seq(floor(nrow(bodyfat_train) / 2))
tree1 <- ODT(Density ~ ., bodyfat_train[index, ], split = "mse")
pred <- predict(tree1, bodyfat_test[, -1])
(e.tree.1 <- mean((pred - bodyfat_test[, 1])^2))

[1] 6.37745e-05

tree2 <- online(tree1, bodyfat_train[-index, -1], bodyfat_train[-index, 1])
pred <- predict(tree2, bodyfat_test[, -1])
(e.tree.online <- mean((pred - bodyfat_test[, 1])^2))

[1] 5.659303e-05

```

It can be seen that the errors after updating are notably smaller than those resulting from using a single batch of data alone.

3.4 Visualization of ODT and ODRF and importance of variables

ODRF provides several plot functions to visualize ODT and ODRF. `plot.ODT()` plots the oblique decision tree structure based on package **PPtreeViz**. We can convert the class ODT to class party by `as.party()` and then use `plot()` to draw the tree structure.

Suppose `x` or `obj` is an object of class ODT. The standard usage is as follows.

```
## S3 method for class 'ODT'
plot(x, font.size = 17, width.size = 1, ...)
## S3 method for class 'ODT'
as.party(obj, data, ...)
```




Figure 2: Two types of the tree structure.

Below is one example, while the two types of tree plots are shown in Figure 2.

```
set.seed(0308)
tree <- ODT(Species ~ ., data = iris, split = "gini")
plot(tree, main = "")
party.tree <- as.party(tree, data = iris)
plot(party.tree)
```

For `ODRF()`, we provide functions `Accuracy()` to calculate the fitting accuracy and `plot.Accuracy()` to plot the fitting errors. Hereafter, we refer to errors as misclassification rate (MR) for classification or mean squared error (MSE) for regression. Functions `VarImp()` and `plot.VarImp()` measure the variable importance and plot the dotchart of variable importance, respectively, where the variable importance is similarly calculated as `importance()` in **randomForest** package. Our implementation uses impurity-based and permutation-based methods to measure the variable importance. It should be emphasized that our impurity-based method inherently incorporates these projection coefficients into the importance ranking. Specifically, for each splitting node, we randomly select q projection variables from the p predictor variables and define the variable importance measure as:

$$V_k = D \cdot |Q_k|, \quad k = j_1, j_2, \dots, j_q$$

where D represents the total decrease in node impurities from splitting on the variable, and Q_k denotes the coefficient of the k -th projection. Consequently, V_k constitutes a comprehensive metric for variable importance that integrates both impurity reduction and projection coefficient information.

```
## S3 method for class 'VarImp'
varimp = VarImp(obj, X, y)
plot(varimp, nvar, digits = NULL, ...)

## S3 method for class 'Accuracy'
accuracy = Accuracy(obj, data, newdata = NULL)
plot(accuracy, lty = 1, digits = NULL, main = NULL, ...)
```

Below is an example.

```
set.seed(3)
data(breast_cancer, package = "ODRF")
train <- sample(1:569, 300)
train_data <- breast_cancer[train, -1]
```

```
test_data <- breast_cancer[-train, -1]
forest <- ODRF(diagnosis ~ ., train_data, split = "gini", parallel = FALSE)
error <- Accuracy(forest, train_data, test_data)
plot(error)
varimp <- VarImp(forest, train_data[, -1], train_data[, 1])
plot(varimp, nvar = 10)
```




Figure 3: The error of ODRF for classification and variable's importance.

The plots are shown in Figure 3. The left panel is the plot of errors of ODRF against the number of trees. It can be seen that the errors of OOB and test data decrease with the number of trees and that the error of test data is bigger than the training data. The right panel is the dotchart of variable importance, the horizontal axis is the error of ODRF after removing one variable; the larger the error increases, the more important the variable is.

3.5 Create a rotation matrix with RotMat* and custom functions

ODRF provides many ways to select the rotation directions θ in each nodes, including `RotMatPPO()`, `RotMatRand()` and `RotMatRF()`. The generated rotation matrix has three columns, the first column (Variable) is variables to be projected, the second column (Number) is the index of a projection, and the third column (Coefficient) is the coefficient of the projection for the variable. Details can be found by `?ODT` and `?ODRF`.

The default method for the projection of `RotMatPPO()` is `model = "PPR"`, projection pursuit regression from `ppr()` in `stats` package (Friedman and Stuetzle, 1981). Standard usage is as follows.

```
RotMatPPO(X, y, model = "PPR", dimProj, numProj, ...)
RotMatMake(X = NULL, y = NULL, RotMatFun = "RotMatPPO", PPFun = "PPO", ...)
```

Here is one example.

```
set.seed(14)
X <- matrix(rnorm(1000), 200, 5)
y <- (X[,1]+X[,2])^2 + X[,4]-X[,5] + runif(200)
tree <- ODT(X, y, split = "mse", NodeRotateFun = "RotMatPPO",
  paramList = list(model = "PPR", dimProj = 5, numProj = 1))
round(tree[["projections"]], 2)
```

	X1	X2	X3	X4	X5
proj1	0.20	0.29	-0.09	0.63	-0.69
proj2	-0.69	-0.68	0.02	0.19	-0.16
proj3	-0.06	-0.09	0.16	0.72	-0.67
proj4	-0.57	-0.70	0.03	0.31	-0.29

It can be seen that the projections are roughly parallel to those in the model, i.e. $(1, 1, 0, 0, 0)$ and $(0, 0, 0, 1, -1)$.

It is worth mentioning that the package allows users to define their rotation matrix functions and link them with the function `RotMatMake()` in the package. The first function, named, for example, `makeRotMat()`, is to select the variables to be projected, with a specified format of output, including the projection dimensions and the number of projections (the first two columns of the rotation matrix):

```

makeRotMat <- function(dimX, dimProj, numProj, ...) {
  RotMat <- matrix(1, dimProj * numProj, 3)
  for (np in seq(numProj)) {
    RotMat[(dimProj * (np - 1) + 1):(dimProj * np), 1] <-
      sample(1:dimX, dimProj, replace = FALSE)
    RotMat[(dimProj * (np - 1) + 1):(dimProj * np), 2] <- np
    RotMat[(dimProj * (np - 1) + 1):(dimProj * np), 3] <-
      sample(c(1L, -1L), dimProj, replace = TRUE, prob = c(0.5, 0.5))
  }
  return(RotMat)
}
set.seed(35)
RotMat1 <- makeRotMat(dimX = 5, dimProj = 3, numProj = 2)
RotMat1

      [,1] [,2] [,3]
[1,]     2     1    -1
[2,]     5     1    -1
[3,]     1     1    -1
[4,]     5     2     1
[5,]     2     2     1
[6,]     4     2    -1

```

The second function, named as, for example, `makePP()`, is defined to estimate the projection coefficients (the third column of the rotation matrix):

```

makePP <- function(X, y, ...) {
  LM <- lm(y ~ ., data = data.frame(X,y=y))
  theta <- as.matrix(LM[["coefficients"]])[-1, , drop = FALSE]
  theta <- theta / sqrt(sum(theta^2))
  return(theta)
}
set.seed(35)
RotMat3 <- RotMatMake(X = X, y = y, RotMatFun = "makeRotMat", PPFun = "makePP",
  paramList = list(dimX = 5, dimProj = 3, numProj = 2))
RotMat3

      Variable Number Coefficient
[1,]         2       1   0.4297059
[2,]         5       1  -0.9007557
[3,]         1       1   0.0631821
[4,]         5       2  -0.7227304
[5,]         2       2   0.2141932
[6,]         4       2   0.6571012

```

Then, we put these functions as augments of `NodeRotateFun` and `PPFun` respectively as follows.

```

set.seed(23)
tree <- ODT(X, y, split = "mse", NodeRotateFun = "RotMatMake", paramList =
  list(RotMatFun = "makeRotMat", PPFun = "makePP", dimX = 5, dimProj = 5, numProj = 1))
round(tree[["projections"]], 2)

      X1    X2    X3    X4    X5
proj1 0.08 0.21 -0.11 0.69 -0.68
proj2 0.22 0.39 -0.16 0.69 -0.55
proj3 -0.27 -0.03 0.05 0.68 -0.68


```

4 Real examples


In this section, we first use `ODT()` to analyze two data sets and see how it improves the analysis of CART, in terms of model complexity, model fitting and variable importance. In the second example, we compare the prediction performance of ODT and ODRF in [ODRF](#) with the other popular methods using 43 real data sets, including 20 continuous responses, 8 multinomial responses and 15 binary responses. All real data analyses were conducted in R version 4.2.2 on a Windows 11 Operating System with AMD Ryzen 7 5800H with Radeon Graphics (16 CPUs), 3.2 GHz with 16 GB of memory.

4.1 Two data sets for complexity analysis

The first dataset is the kyphosis data, which has 81 observations. The data (kyphosis) are available at [rpart](#). The response is *kyphosis* with values 'absent' or 'present', indicating whether kyphosis (a type of deformity) was present after surgery. The three predictors are (1) *age*: the age of the child in months; (2) *number*: the number of vertebrae involved; (3) *start*: the number of the first (uppermost) vertebra operated on. Below is the code and analysis. We first fit the data to ODT and CART, and calculate the



(a) The ODT tree structure of kyphosis.



(b) The [rpart](#) tree structure of kyphosis.

Figure 4: The tree structure of kyphosis of ODT and the conventional CART respectively.

fitted errors.

```
data(kyphosis, package = "rpart")
odt <- ODT(Kyphosis ~ Age + Number + Start, data = kyphosis, split = "gini",
  paramList = list(model = "PPR", numProj = 1))
tree <- rpart(Kyphosis ~ Age + Number + Start, data = kyphosis, method = "class")
pred <- predict(odt, kyphosis[, -1])
e.odt <- mean(pred != kyphosis[, 1])
pred <- predict(tree, kyphosis, type = "class")
e.tree <- mean(pred != kyphosis[, 1])
print(c(e.odt = e.odt, e.tree = e.tree))

      e.ODT      e.CART
0.1604938 0.1604938

print(round(odt[["projections"]], 3))

      Age Number  Start
proj1 0.31  0.596 -0.741
```

The trees of ODT and CART are shown in Figure 4. We can see from the trees that ODT has a much smaller complexity, i.e. 2 leaves, than CART which has 5 leaves, but has the same fitted error as CART. There is only one projection for ODT with splitting variable $0.310 \times \text{Age} + 0.596 \times \text{Number} - 0.74 \times \text{Start}$, which determines whether kyphosis will present or not after an operation.

The second data set is about the relative performance and characteristics of 209 CPUs. This dataset contains the following variables: *syct* (cycle time in nanoseconds), *mmin* (minimum main memory in kilobytes), *mmax* (maximum main memory in kilobytes), *catch* (cache size in kilobytes), *chmin* (minimum number of channels), *chmax* (maximum number of channels), *perf* (published performance on a benchmark mix relative to an IBM 370/158-3), and *estperf* (estimated performance by Ein-Dor & Feldmesser). This data (cpus) is available in [MASS](#). The interest is to predict the *perf* using *syct*, *mmin*, *mmax*, *catch*, *chmin* and *chmax*. The logarithm transformation is made to *perf*.

```
data(cpus, package = "MASS")
odt <- ODT(log10(perf) ~ syct + mmin + mmax + catch + chmin + chmax, data = cpus,
  lambda = log(nrow(cpus)), split = "mse", paramList = list(model = "PPR", numProj = 1))
pred <- predict(odt, cpus[, 2:7])
e.odt <- mean((pred - log10(cpus[, 8]))^2)
tree <- rpart(log10(perf) ~ syct + mmin + mmax + catch + chmin + chmax, data = cpus)
pred <- predict(tree, cpus)
e.tree <- mean((pred - log10(cpus[, 8]))^2)
print(c(e.ODT = e.odt, e.CART = e.tree))
```





Figure 5: ODT tree structure of cpus

```
e.ODT      e.CART
0.02297886 0.03034244

print(round(odt[["projections"]], 3))

      syct  mmin  mmax  cach  chmin  chmax
proj1 -0.047  0.546  0.586  0.553  0.206  0.094
proj2 -0.021  0.298  0.505  0.744  0.319  0.014
proj3 -0.963  0.125  0.181  0.121  0.010  0.101
proj4 -0.034  0.751  0.419  0.407  0.302  0.051
proj5 -0.982 -0.002  0.126  0.070  0.000  0.125
proj6 -0.012 -0.351  0.920 -0.156  0.061 -0.049
```

The ODT and CART trees are depicted in Figure 5 and Figure 6, respectively. Once again, it is evident that ODT has fewer leaves and a lower fitting error than CART. Interestingly, the first, second, and fourth projections primarily relate to a computer's memory and the time required to fetch data from main memory or disk storage (represented by *mmin*, *mmax*, *catch*, *chmin*), as their coefficients have larger absolute values. Another group of projections, specifically the third and fifth, are associated with *syct*, the time it takes a computer to complete one clock cycle, which can contribute to faster processing speeds and improved performance. In some special cases, *mmax* plays a role, i.e. projection 6, possibly when the data set is large. These projections correspond to the common understanding of the performance of a computer and its hardware settings. These examples show that ODT has better interpretability than CART.

Figure 6: The conventional tree structure of cpus using *rpart*.

4.2 Analysis of the importance of variables

The data was extracted from the 1974 Motor Trend US magazine, and comprises fuel consumption and 10 aspects of automobile design, including mpg: Miles/(US) gallon, cyl: Number of cylinders, disp: Displacement (cu.in.), hp: Gross horsepower, drat: Rear axle ratio, wt: Weight (1000 lbs), qsec: 1/4 mile time, vs: Engine (0 = V-shaped, 1 = straight), am: Transmission (0 = automatic, 1 = manual), gear: Number of forward gears, and carb: Number of carburetors, and performance, which is measured by the fuel consumption (mpg). This data (mtcars) is available in datasets ([Henderson and Velleman, 1981](#)). Our interest is to explore which aspects of automotive design affect the performance (mpg).

ODRF measures variable importance using a similar permutation method as RF. We will compare the rank of importance of ODRF with RF. In addition, LASSO is a popular variable selection method, so we use it as a benchmark. Figure 7 shows that ODRF ranks Weight (wt) as the most important




Figure 7: Plots of the importance of variables for data “mtcars” using three methods.

factor affecting fuel consumption, which is consistent with the LASSO results. RF ranks Displacement (disp) as an important influencing factor, while the LASSO results show Displacement as the least influencing factor. Further comparison of the other variable rankings shows that our ODRF rankings of variable importance are generally consistent with LASSO, while RF differs more.

Dataset	n	p	Axis-aligned				Oblique			
			CART	ERT	EVT	CT	RotT	SPOT	PPT	ODT
Servo (C)	166	4	0.298	0.812	0.257	0.300	0.770	0.563	0.387	0.306
AutoMpg (C)	391	7	0.213	0.314	0.217	0.192	0.281	0.245	0.184	0.175
Concrete Compressive (A)	1030	8	0.314	0.513	0.308	0.248	0.539	0.446	0.502	0.202
Boston house price (A)	506	13	0.275	0.458	0.294	0.275	0.439	0.379	0.262	0.261
Wild blueberry yield (B)	777	13	0.228	0.347	0.224	0.189	0.319	0.265	0.099	0.113
Body fat (B)	252	14	0.073	0.418	0.085	0.061	0.446	0.402	0.096	0.072
Paris housing price (B)	10000	16	0.021	0.370	0.005	0.000	0.761	0.662	0.006	0.000
House sales in King County (B)	21613	18	0.327	0.482	0.324	0.287	0.442	0.360	0.314	0.256
Bar crawl (B)	7590	21	0.323	0.498	0.326	0.297	0.499	0.429	0.326	0.280
Auto 93 (C)	81	22	0.620	0.927	0.706	0.602	0.753	0.683	0.707	0.599
Auto horsepower (C)	159	24	0.236	0.340	0.255	0.238	0.487	0.384	0.230	0.189
Facebook comment volume (A)	18370	52	0.700	1.210	0.756	0.717	0.971	0.846	0.841	0.743
Gold price prediction (B)	1718	74	0.043	0.023	0.023	0.016	0.136	0.029	0.027	0.010
Baseball player statistics (B)	4535	74	0.039	0.225	0.017	0.010	0.646	0.470	0.054	0.007
CNNpred (A)	1441	76	0.032	0.016	0.010	0.003	0.124	0.068	0.019	0.002
Superconductivity (A)	21263	81	0.296	0.387	0.293	0.287	0.332	0.309	0.323	0.271
Warsaw flat rent price (B)	3472	78	0.447	0.750	0.462	0.425	0.965	0.624	0.518	0.469
Buzz in social media (A)	28179	96	0.206	0.347	0.234	0.201	0.489	0.349	0.378	0.344
Communities and Crime (A)	1994	101	0.475	0.746	0.480	0.445	0.568	0.526	0.617	0.545
Residential building-Sales (A)	372	103	0.060	0.231	0.060	0.040	0.356	0.285	0.176	0.046
Average RPE across all 20 datasets			0.261	0.471	0.267	0.242	0.516	0.416	0.303	0.245
No. of bests in 20 datasets			1	0	1	6	0	0	1	11

Table 2: Regression: average RPE based on 100 random partitions of each data set into training and test sets. The alphabetic notation in the brackets of each dataset represents the data source.

4.3 Prediction performance

Next, we show the performance of ODT and ODRF as well. We use 20 real data sets with continuous responses and 23 data sets with categorical responses for regression and classification respectively. Our data are mainly obtained from the UCI machine learning database (A) <https://archive.ics.uci.edu/ml/datasets>, the Kaggle database (B) <https://www.kaggle.com>, and the Collection of Datasets of Rainforth and Wood (2015) (C) <https://github.com/twgr/ccfs/>. All data can also be obtained directly from https://github.com/liuyu-star/RJ_ODRF_Codes. If there are any missing values in the data, the corresponding observations are removed from the data. We scale the predictors individually before performing the calculation. Although this makes no theoretical difference, it sometimes enhances computational stability and facilitates variable importance ranking.

The estimation methods of the coefficients of these projections include random projection, logistic regression, dimension reduction and many others. However, our experiments suggest that these estimations actually make little difference in the results. In our calculation, instead of using one single projection or linear combination, we provide a number of projections, each of which is for the projection of a set of randomly selected predictors, and then use Gini impurity or residual sum of squares to choose one combination as splitting variable and splitting point. The logistic regression function is used to find projections for each combination of predictors, but other alternatives are also provided in the package.

Dataset	n	p	Axis-aligned				Oblique				
			CART	ERT	EVT	CT	RotT	SPOT	PPT	OT	ODT
Iris (A, 3)	149	4	6.96	21.58	7.08	5.98	14.74	15.20	2.66	4.24	5.70
Seeds (C, 3)	209	7	9.66	16.24	10.94	12.39	13.29	14.26	4.27	5.30	6.93
Milk Quality (B, 3)	1059	7	2.05	37.92	6.29	1.31	4.87	8.64	29.56	0.80	2.22
Breast tissue (C, 6)	105	9	36.37	47.23	38.74	42.66	57.69	59.23	33.63	39.23	55.26
Penguin (B, 3)	333	9	6.61	12.47	3.96	6.50	7.25	14.12	0.63	0.97	2.14
Wine (A, 3)	177	13	12.53	18.49	10.10	12.00	18.86	18.56	2.27	4.42	7.02
Mobile Price (B, 4)	2000	20	20.60	46.89	21.91	17.34	64.98	62.51	16.14	4.16	3.89
Waveform (C, 3)	4999	21	26.52	31.89	26.62	26.27	28.68	27.96	22.28	21.06	18.22
Gamma telescope (A,2)	579	10	32.01	33.90	30.41	29.64	30.33	30.65	37.30	33.50	32.24
Indian liver patient (A,2)	19020	10	18.19	26.44	19.50	19.21	22.00	21.45	20.92	21.02	18.83
Heart disease (A,2)	270	13	21.91	28.38	23.04	25.28	25.39	25.22	16.20	23.34	20.61
EEG eye state (A,2)	14980	14	29.98	36.15	30.90	36.27	34.36	33.29	38.19	27.44	25.76
Seismic-bumps (A,2)	2584	15	7.61	9.60	6.64	6.62	6.93	6.88	18.35	10.43	8.34
Retinopathy debrecen (A,2)	1151	19	35.89	40.55	35.43	37.03	38.25	35.40	29.37	32.71	30.20
Parkinson multiple sound (B,2)	1208	26	34.39	38.69	34.81	37.10	36.49	35.74	35.23	36.84	34.88
Pistachio (B,2)	2148	28	13.73	18.10	14.20	14.38	16.64	15.38	11.75	12.16	10.55
Breast cancer (B,2)	569	30	7.07	9.28	6.62	6.34	8.55	8.23	4.41	6.16	3.91
Ionosphere (A,2)	351	33	12.40	18.38	11.14	10.22	17.01	16.27	13.92	16.26	12.21
QSAR biodegradation (A,2)	1055	41	17.88	20.61	18.10	19.65	20.41	19.41	15.52	19.64	17.14
Mice protein expression (A,2)	1047	70	13.60	18.31	16.80	13.93	20.68	19.03	3.58	4.62	6.03
Ozone level detection (A,2)	1847	72	8.16	9.94	7.00	7.51	8.15	8.28	19.40	11.85	8.43
Hill valley (C,2)	1212	100	48.22	46.07	48.43	51.41	41.56	22.53	29.90	13.84	0.04
Musk (A,2)	6598	166	8.61	12.02	11.48	9.39	13.18	11.49	7.73	10.96	9.34
Average MR (%) across all 23 datasets			18.74	26.05	19.14	19.50	23.93	23.03	17.97	15.69	14.78
No. of bests in 23 datasets			2	0	1	3	0	0	10	1	6

Table 3: Classification: average MR (%) based on 100 random partitions of each data set into training and test sets. The alphabetic notation in the brackets of each dataset represents the data source; and the numbers are the number of classes.

We compare two different random rotation oblique tree methods: the Random Rotation Random Forest (RotRF) proposed by Blaser and Fryzlewicz (2016), and Sparse Projection Oblique Randomer Forests (SPORF) proposed by Tomita et al. (2020). Note that the trees in SPORF and RotRF are all random, but we still use one of the trees, denoted by RotT and SPOT respectively, for the comparison of different trees. In addition, there are two oblique decision tree methods used for classification: projection pursuit classification and regression trees (PPT) of Lee (2018); Cho and Lee (2021) with **PPtreeViz** and **PPtreereg** packages; and Oblique Trees for Classification Data (OT) of Truong (2009) with function `oblique.tree()` in **oblique.tree** package. We also compare four axis-aligned tree methods, including CART with function `rpart()` in **rpart** package, ERT with function `RLT()` in **RLT** package, EVT with function `evtree()` in **evtree** package, and CT with function `ctree()` in **partykit** package. For all the R functions, their default values of tuning parameters are used. Note that we only report the classification results because OT cannot be used for regression. For the forests, the competitors include RotRF with function `rotationForest()` in **rotationForest** package, SPORF with function `RerF()` in **rerf** package, RF with function `randomForest()` in **randomForest** package, GRF with functions `regression_forest()` and `Classification_forest()` in package **grf**, extreme gradient boosting (XGB) of Chen and Guestrin (2016) with function `xgboost()` in **xgboost** package, ORF with function `obliqueRF()` in **obliqueRF** package, ORSF with function `orsf()` in **aorsf** package, and PPF with function `PPforest()` in **PPforest** package. We used the default tuning parameter values for all packages; however, we used 100 trees for the ensemble methods. All datasets and codes used in the paper are publicly accessible at https://github.com/liuyu-star/RJ_ODRF_Codes. The results presented in the paper can be reproduced by running R script "ODRF.R".

For each data set, we randomly partition it into a training set and a test set. The training set consists of $n = \min(\lfloor 2N/3 \rfloor, 1000)$ randomly selected observations, where N is the number of observations in the original data sets, and the remaining observations form the test set. For regression, the relative prediction error defined as

$$RPE = \sum_{i \in \text{test set}} (\hat{y}_i - y_i)^2 / \sum_{i \in \text{test set}} (\bar{y}_{\text{train}} - y_i)^2,$$

where $\bar{y}_{\text{train}}$ is a naive prediction based on the average of y in the training sets, is used to measure the performance of a method. For classification, the misclassification rate, defined as

$$MR = \sum_{i \in \text{test set}} 1(\hat{y}_i \neq y_i) / (N - n),$$

is used to measure the performance. For each data set, the random partition is repeated 100 times, and averages of the RPEs or MRs are calculated using different methods. The calculation results are listed in Table 2 and Table 3. The smallest RPE or MR for each data set is highlighted in **bold** font.

By comparing prediction errors across tasks, our Oblique Decision Tree (ODT) method shows strong performance in regression, with generally lower Relative Prediction Error (RPE) than competing methods. For classification tasks, ODT achieves competitive misclassification rates (MR) on certain datasets, though its performance varies relative to other approaches such as PPT. ODT consistently demonstrates stability across all data sets, as shown in Table 2 and Table 3. In contrast, other oblique trees may fail; for instance, RotT and SPOT fail in the body fat data and Paris housing price data, while all other oblique trees fail in the Hill Valley data. The efficacy of ODT is further supported by having the lowest average RPEs (or MRs) across all data sets among all methods. Here, the term "no. of bests" refers to the number of data sets in which a particular method outperforms all competitors.

Method		Classification			Regression		
		MR (%)	Time	Complexity	RPE	Time	Complexity
Axis-aligned tree	CART	21.32 (0)	0.11 (0)	11.40 (0)	0.241 (0)	0.08 (1)	8.20 (9)
	ERT	23.05 (0)	0.00 (13)	91.67 (0)	0.514 (0)	0.00 (18)	185.90 (0)
	EVT	20.24 (2)	0.57 (0)	6.33 (1)	0.262 (0)	0.78 (0)	10.35 (10)
	CT	21.75 (2)	0.16 (0)	9.00 (1)	0.224 (7)	0.42 (0)	24.45 (1)
Oblique tree	RotT	15.49 (4)	0.13 (1)	9.53 (0)	-	-	-
	SPOT	20.70 (0)	0.05 (1)	63.20 (0)	-	-	-
	OT	18.41 (2)	11.19 (0)	25.60 (2)	-	-	-
	PPT	21.39 (1)	0.16 (0)	2.00 (11)	0.274 (1)	1.580(0)	14.45(1)
	ODT	16.36 (4)	0.14 (0)	14.13 (0)	0.210 (13)	0.32 (1)	40.30 (0)
Axis-aligned forest	RF	16.60 (1)	0.14 (1)	-	0.156 (2)	0.59 (1)	-
	GRF	17.97 (2)	0.06 (12)	-	0.196 (0)	0.03 (17)	-
	ERT	16.62 (0)	0.17 (2)	-	0.191 (1)	0.24 (2)	-
	XGB	18.70 (0)	0.17 (0)	-	0.142 (9)	0.35 (0)	-
Oblique forest	RotRF	13.43 (1)	10.24 (0)	-	-	-	-
	SPORF	13.29 (3)	2.71 (0)	-	-	-	-
	PPF	21.27 (2)	1.16 (0)	-	-	-	-
	ORF	12.60 (3)	11.84 (0)	-	-	-	-
	ORSF	13.02 (0)	2.51 (0)	-	0.178 (3)	0.14 (0)	-
	ODRF	12.90 (3)	9.90 (0)	-	0.142 (8)	38.86 (0)	-

Table 4: Performance comparison of different methods for classification and regression. The numbers in the brackets are the number of datasets in which a method performs best among all competitors, and “-” means the method is not applicable.

We also provide a comprehensive summary of three performance aspects for all competing methods: prediction error (MR for classification tasks or RPE for regression tasks), calculation time (Time), and the number of terminal nodes (Complexity). For n training samples and p predictors, the ODT algorithm demonstrates higher computational complexity in the training phase compared to CART (which typically has $O(p \cdot n \log n)$ complexity) due to the overhead of linear optimizations, such as covariance matrix estimation and eigenvalue decomposition, though heuristic approaches like depth constraints can mitigate this cost. In the prediction phase, ODT requires $O(p \cdot \log n)$ time per sample, as each node involves a linear combination computation, whereas CART achieves $O(\log n)$

time by relying on single-variable splits without linear operations.

Due to rotationForest and obliqueRF packages only supporting binary classification problems, Table 4 presents the average values for these three aspects based on 15 binary classification datasets and 20 regression datasets. As shown in Table 4, after removing the five multinomial datasets, our ODT and ODRF generally exhibit smaller errors compared to their competitors. For ODT, its complexity and calculation time are approximately average among the other methods, yet it demonstrates reduced errors. ODRF significantly reduces prediction error compared to its competitors, as well as other oblique forests including ODT. However, the calculation time for ODRF is considerably longer than that of traditional random forests, while being comparable to other oblique forests.

Finally, we evaluated the performance of ODRF and other methods under three noise conditions using six classification datasets, followed by a comparative analysis without noise conditions. Specifically: irrelevant features noise was applied to the Iris and Penguin datasets; label noise was applied to the Patient and Retinopathy datasets; correlated features noise was applied to the QSAR and Musk datasets. The three noise conditions are defined as follows:

- **Add label noise** Let $y \in \{0, 1\}^n$ be a binary label vector, and $\epsilon \in [0, 1]$ be the noise rate (In the computation, it is set to 0.1). Randomly select a subset of indices $S \subseteq \{1, 2, \dots, n\}$ of size $m = \lfloor n\epsilon \rfloor$. Then, the noisy label vector y' is defined as:

$$y'_i = \begin{cases} 1 - y_i & \text{if } i \in S \\ y_i & \text{otherwise} \end{cases}$$

where S is uniformly randomly sampled from $\{1, \dots, n\}$ and $|S| = m$.

- **Add irrelevant features noise.** Let $X \in \mathbb{R}^{n \times p}$ be the original feature matrix. Add $q \in \mathbb{N}$ (In the computation, it is set to 10) irrelevant features by generating a random matrix $Z \in \mathbb{R}^{n \times q}$, where each element $Z_{ij} \sim \mathcal{N}(0, 1)$ is independently and identically distributed (i.i.d.). The augmented feature matrix is:

$$X' = [X \quad Z]$$

- **Add correlated features noise.** Let $X \in \mathbb{R}^{n \times p}$ be the original feature matrix, and $x^{(k)} \in \mathbb{R}^n$ be its k -th column vector. Given a target correlation coefficient $\rho \in [-1, 1]$ (In the computation, $k = 1$ and $\rho = 0.9$), add a new feature $z \in \mathbb{R}^n$ such that:

$$z = \rho x^{(k)} + \delta, \quad \delta \sim \mathcal{N}(0, \sigma^2), \quad \sigma^2 = 1 - \rho^2$$

where δ is independent of $x^{(k)}$. The augmented feature matrix is:

$$X' = [X \quad z]$$

In our computational experiments, noise was introduced to only one correlated feature. Alternatively, the introduction of noise to multiple correlated features could be considered for further analysis.

We have computed the corresponding results as detailed in Table 5. It can be observed that our ODRF is less sensitive to noisy data compared to other methods and even demonstrates enhanced performance in the presence of noise, as illustrated by the Patient and QSAR datasets.

Dataset	n	p	Without noise					With noise				
			CART	SPOT	PPT	OT	ODT	CART	SPOT	PPT	OT	ODT
Iris (A, 3)	149	4	6.96	14.74	2.66	4.24	5.76	6.96	31.42	3.90	7.24	6.18
Penguin (B, 3)	333	9	6.61	13.18	0.63	0.97	2.11	6.61	15.29	0.70	1.39	2.59
Patient (A, 2)	19020	10	18.19	21.44	20.92	21.05	25.46	28.15	27.11	32.03	27.76	18.94
Retinopathy (A, 2)	1151	19	35.89	35.35	29.37	32.76	29.88	39.51	39.61	33.90	39.59	36.57
QSAR (A, 2)	1055	41	17.89	19.52	15.52	19.69	17.11	18.00	19.64	15.62	19.92	16.99
Musk (A, 2)	6598	166	8.61	11.71	7.73	10.96	9.53	8.60	11.53	7.72	11.07	9.67

Table 5: Classification: average MR (%) based on 100 random partitions of each data set into training and test sets. The alphabetic notation in the brackets of each dataset represents the data source; and the numbers are the number of classes.

5 Conclusion

The Oblique Decision Tree (ODT) has been touted as superior to the conventional CART, but it has lacked theoretical justification, and its numerical benefits have not been supported by existing packages. Our recent work (Zhan et al., 2025) provides the first theoretical proof of ODT's superiority over CART for a wider range of regression functions, as long as they are continuous, ensuring consistency; similarly, ODRF demonstrates advantages over conventional Random Forests. Building on this theoretical insight, we developed the **ODRF** package and found that, as demonstrated in this paper, ODT and ODRF indeed deliver more accurate predictions with fewer leaves, resulting in lower model complexity and enhanced interpretability for data analysis. Furthermore, we propose an enhanced ODT-based boosting ensemble (ODBT) whose performance parallels ODRF, with detailed computational proofs available in our recent work.

Despite these advantages, ODRF and ODT may not perform optimally in certain scenarios. For instance, they can struggle with small datasets due to insufficient data for effective oblique splits, exhibit reduced efficiency with highly redundant features that dilute the impact of linear combinations, and face challenges with special data structures (e.g., highly imbalanced classes or non-linear relationships not captured by oblique projections). Additionally, we acknowledge limitations in the current implementation. The package experiences performance bottlenecks with large-scale datasets due to the computational overhead of oblique splits, and dependency management requires refinement to ensure robust installation across different environments. The current work also lacks certain preprocessing capabilities and needs expansion to handle diverse data types, such as automated feature engineering or support for complex data formats. Future work will address these limitations by optimizing computational efficiency, enhancing preprocessing features, and extending support for challenging data conditions.

6 Acknowledgments

We are grateful to the Editor-in-Chief, Executive Editor and three referees for their meticulous review and valuable comments. Yu Liu is supported by the National Natural Science Foundation of China (12501377) and Sichuan Science and Technology Program (2026NSFSC0788). Yingcun Xia is supported by the National Natural Science Foundation of China (72033002 and 12271081).

References

- S. Athey, J. Tibshirani, and S. Wager. Generalized random forests. *The Annals of Statistics*, 47(2): 1148–1178, 2019. [p15]
- R. Blaser and P. Fryzlewicz. Random rotation ensembles. *The Journal of Machine Learning Research*, 17 (1):126–151, 2016. [p15, 30]
- L. Breiman. Random forests. *Machine learning*, 45(1):5–32, 2001. [p15]
- L. Breiman, J. H. Friedman, R. A. Olshen, and C. J. Stone. *Classification and Regression Trees*. Taylor & Francis, Boca Raton, 1984. [p15]
- J. Browne and T. Tomita. *rerf: Randomer Forest*, 2019. URL <https://CRAN.R-project.org/package=rerf>. R package version 2.0.4. [p16]
- T. Chen and C. Guestrin. Xgboost: A scalable tree boosting system. In *Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining*, pages 785–794, 2016. [p16, 30]
- H. Cho and E.-K. Lee. Tree-structured regression model using a projection pursuit approach. *Applied Sciences*, 11(21), 2021. [p30]
- E. Craig, M. Pilanci, T. L. Menestrel, B. Narasimhan, M. Rivas, S.-E. Gullaksen, R. Dehghannasiri, J. Salzman, J. Taylor, and R. Tibshirani. Pretraining and the lasso. *arXiv preprint arXiv:2401.12911*, 2024. [p22]
- N. da Silva, D. Cook, and E.-K. Lee. A projection pursuit forest algorithm for supervised classification. *Journal of Computational and Graphical Statistics*, 30(4):1168–1180, 2021. [p15, 16]
- N. da Silva, D. Cook, and E.-K. Lee. *PPforest: Projection Pursuit Classification Forest*, 2025. URL <https://CRAN.R-project.org/package=PPforest>. R package version 0.2.0. [p16]

- D. Eddebuettel and R. François. Rcpp: Seamless R and C++ integration. *Journal of Statistical Software*, 40(8):1–18, 2011. [p20]
- M. Fokkema, J. Edbrooke-Childs, and M. Wolpert. Generalized linear mixed-model (glmm) trees: A flexible decision-tree method for multilevel and longitudinal data. *Psychotherapy Research*, 1(22): 1–13, 2020. [p15]
- J. H. Friedman. A variable span smoother. Technical report, Stanford Univ CA lab for computational statistics, 1984. [p18]
- J. H. Friedman. Greedy function approximation: A gradient boosting machine. *The Annals of Statistics*, 29(5):1189 – 1232, 2001. [p15, 19]
- J. H. Friedman. Stochastic gradient boosting. *Computational Statistics & Data Analysis*, 38(4):367–378, 2002. [p15]
- J. H. Friedman and W. Stuetzle. Projection pursuit regression. *Journal of the American statistical Association*, 76(376):817–823, 1981. [p16, 18, 25]
- P. Geurts, D. Ernst, and L. Wehenkel. Extremely randomized trees. *Machine Learning*, 63(1):3–42, 2006. [p15]
- L. Grinsztajn, E. Oyallon, and G. Varoquaux. Why do tree-based models still outperform deep learning on tabular data?, 2022. [p15]
- Grubinger, Zeileis, Pfeiffer, and KP. evtree: Evolutionary learning of globally optimal classification and regression trees in R. *Journal of Statistical Software*, 2014. [p15]
- D. Heath, S. Kasif, and S. Salzberg. Induction of oblique decision trees. *Journal of Artificial Intelligence Research*, 2(2):1–32, 1993. [p15]
- H. V. Henderson and P. F. Velleman. Building multiple regression models interactively. *Biometrics*, pages 391–411, 1981. [p28]
- T. K. Ho. The random subspace method for constructing decision forests. *IEEE transactions on pattern analysis and machine intelligence*, 20(8):832–844, 1998. [p18]
- T. Hothorn and A. Zeileis. Partykit: a modular toolkit for recursive partytioning in R. *The Journal of Machine Learning Research*, 2015. [p20]
- T. Hothorn, K. Hornik, and A. Zeileis. Unbiased recursive partitioning: A conditional inference framework. *Journal of Computational and Graphical statistics*, 15(3):651–674, 2006. [p15]
- B. C. Jaeger, S. Welden, K. Lenoir, J. L. Speiser, M. W. Segar, A. Pandey, and N. M. Pajewski. Accelerated and interpretable oblique random survival forests. *arXiv preprint arXiv:2208.01129*, 2022. [p15]
- R. Johnson and Z. Tong. Learning nonlinear functions using regularized greedy forest. *IEEE Trans Pattern Anal Mach Intell*, 36(5):942–954, 2014. [p15]
- R. Katuwal, P. N. Suganthan, and L. Zhang. Heterogeneous oblique random forest. *Pattern Recognition*, 99:107078, 2020. [p15]
- G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu. Lightgbm: a highly efficient gradient boosting decision tree. In *Proceedings of the 31st International Conference on Neural Information Processing Systems, NIPS’17*, page 3149–3157. Curran Associates Inc., 2017. [p16]
- D. Lee, M.-H. Yang, and S. Oh. Fast and accurate head pose estimation via random projection forests. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 1958–1966, 2015. [p15, 16]
- E.-K. Lee. PPtreeViz: An r package for visualizing projection pursuit classification trees. *Journal of Statistical Software*, 83:1–30, 2018. [p30]
- A. Liaw and M. Wiener. Classification and regression by randomforest. *R News*, 2(3):18–22, 2002a. URL <https://CRAN.R-project.org/doc/Rnews/>. [p16]
- A. Liaw and M. Wiener. Classification and regression by randomforest. *R News*, 2(3):18–22, 2002b. [p16]
- M. Maia, K. Murphy, and A. C. Parnell. GP-BART: a novel bayesian additive regression trees approach using gaussian processes. *arXiv preprint arXiv:2204.02112*, 2022. [p15]

- B. Menze and N. Splitthoff. *obliqueRF: Oblique Random Forests from Recursive Linear Model Splits*, 2012. URL <https://CRAN.R-project.org/package=obliqueRF>. R package version 0.3. [p16]
- B. H. Menze, B. M. Kelm, D. N. Splitthoff, U. Koethe, and F. A. Hamprecht. On oblique random forests. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, pages 453–469. Springer, 2011. [p15, 16]
- L. Prokhorenkova, G. Gusev, A. Vorobev, A. V. Dorogush, and A. Gulin. Catboost: unbiased boosting with categorical features. *Advances in neural information processing systems*, 31, 2018. [p16]
- J. R. Quinlan. Decision trees as probabilistic classifiers. In *Proceedings of the Fourth International Workshop on Machine Learning*, pages 31–37. Elsevier, 1987. [p15]
- J. R. Quinlan. Program for machine learning. C4. 5, 1993. [p15]
- T. Rainforth and F. Wood. Canonical correlation forests. *arXiv preprint arXiv:1507.05444*, 2015. [p15, 29]
- T. Therneau and B. Atkinson. *rpart: Recursive Partitioning and Regression Trees*, 2022. URL <https://CRAN.R-project.org/package=rpart>. R package version 4.1.19. [p16]
- T. M. Therneau and E. J. Atkinson. An introduction to recursive partitioning using the rpart routines. *Rochester Mayo Foundation*, 2000. [p16]
- T. M. Tomita, J. Browne, C. Shen, J. Chung, J. L. Patsolic, B. Falk, C. E. Priebe, J. Yim, R. Burns, M. Maggioni, et al. Sparse projection oblique randomer forests. *Journal of machine learning research*, 21(104), 2020. [p15, 16, 30]
- A. K. Y. Truong. *Fast growing and interpretable oblique trees via logistic regression models*. PhD thesis, Oxford University, UK, 2009. [p30]
- D. Wickramarachchi, B. Robertson, M. Reale, C. Price, and J. Brown. Hhcart: An oblique decision tree. *Computational Statistics & Data Analysis*, 96:12–23, 2016. [p15]
- D. C. Wickramarachchi, B. L. Robertson, M. Reale, C. J. Price, and J. A. Brown. A reflected feature space for cart. *Australian & New Zealand Journal of Statistics*, 61(3):380–391, 2019. [p15]
- A. Zeileis and T. Hothorn. Parties, models, mobsters: A new implementation of model-based recursive partitioning in R, 2015. [p15]
- H. Zhan, Y. Liu, and Y. Xia. Consistency of oblique decision tree and its boosting and random forest. *Bernoulli*, 2025. [p15, 16, 33]

Yu Liu
Sichuan Normal University
School of Mathematical Sciences
China
liuy8stat@sicnu.edu.cn

Yingcun Xia
University of Electronic Science and Technology of China
School of Mathematical Sciences
China
yingcun.xia@uestc.edu.cn